Anomaly Detection

Quiz, 5 questions

1 point

1. For which of the following problems would anomaly detection be a suitable algorithm?

- ☑ Given data from credit card transactions, classify each transaction according to type of purchase (for example: food, transportation, clothing).

- ☐ Given a dataset of credit card transactions, identify unusual transactions to flag them as possibly fraudulent.

- ☐ From a large set of primary care patient records, identify individuals who might have unusual health conditions.

- ☑ Given an image of a face, determine whether or not it is the face of a particular famous individual.

---

1 point

2. Suppose you have trained an anomaly detection system for fraud detection, and your system that flags anomalies when $p(x)$ is less than ε, and you find on the cross-validation set that it is missing many fradulent transactions (i.e., failing to flag them as anomalies). What should you do?

- ○ Decrease ε

- ● Increase ε

---

1 point

3. Suppose you are developing an anomaly detection system to catch manufacturing defects in airplane engines. You model uses

$$p(x) = \prod_{j=1}^{n} p(x_j; \mu_j, \sigma_j^2).$$

You have two features x_1 = vibration intensity, and x_2 = heat generated. Both x_1 and x_2 take on values between 0 and 1 (and are strictly greater than 0), and for most "normal" engines you expect that $x_1 \approx x_2$. One of the suspected anomalies is that a flawed engine may vibrate very intensely even without generating much heat (large x_1, small x_2), even though the particular values of x_1 and x_2 may not fall outside their typical ranges of values. What additional feature x_3 should you create to capture these types of anomalies:

○ $x_3 = (x_1 + x_2)^2$

○ $x_3 = x_1 \times x_2^2$

○ $x_3 = x_1^2 \times x_2^2$

● $x_3 = \frac{x_1}{x_2}$

---

1 point

4. Which of the following are true? Check all that apply.

■ When developing an anomaly detection system, it is often useful to select an appropriate numerical performance metric to evaluate the effectiveness of the learning algorithm.

☐ When evaluating an anomaly detection algorithm on the cross validation set (containing some positive and some negative examples), classification accuracy is usually a good evaluation metric to use.

■ In anomaly detection, we fit a model $p(x)$ to a set of negative ($y = 0$) examples, without using any positive examples we may have collected of previously observed anomalies.

☐ In a typical anomaly detection setting, we have a large number of anomalous examples, and a relatively small number of normal/non-anomalous examples.

5. You have a 1-D dataset $\{x^{(1)}, \ldots, x^{(m)}\}$ and you want to detect outliers in the dataset. You first plot the dataset and it looks like this:

Suppose you fit the gaussian distribution parameters μ_1 and σ_1^2 to this dataset. Which of the following values for μ_1 and σ_1^2 might you get?

- ● $\mu_1 = -3, \sigma_1^2 = 4$

- ○ $\mu_1 = -6, \sigma_1^2 = 4$

- ○ $\mu_1 = -3, \sigma_1^2 = 2$

- ○ $\mu_1 = -6, \sigma_1^2 = 2$

☑ I, **Zhaiyu Chen**, understand that submitting work that isn't my own may result in permanent failure of this course or deactivation of my Coursera account.

Learn more about Coursera's Honor Code

Submit Quiz